

O&M runbooks — steady-state operations & maintenance

Operations and maintenance for a **deployed** Claude apps gateway. The [test-run-runbook](#) covers the initial deploy; this document covers what you do afterwards: rotating the certificate and secrets, refreshing the RDS CA bundle, pushing new Claude Code / image versions, responding to alarms, backup/restore, and teardown.

Every command uses the repo's own scripts and `deploy.env` variables — never hardcoded org values. Run operator commands from a host that has `deploy.env` filled in and AWS credentials for `us-gov-west-1` (the scripts source `scripts/common.sh`, which loads `deploy.env` and exports `AWS_REGION`). The three stack names come from `deploy.env`:

- `DB_STACK_NAME` — `01-database.yaml` (default `${NAME_PREFIX}-db`)
- `GATEWAY_STACK_NAME` — `02-gateway.yaml` (default `${NAME_PREFIX}`)
- `OBS_STACK_NAME` — `03-observability.yaml` (default `${NAME_PREFIX}-obs`)

Verification status — read this first

Per the repo honesty rule (`.claude/rules/process.md`), each runbook is tagged by how far its steps have been exercised on real infrastructure. The source of truth is the fix log at the top of [security-review-2026-07.md](#) and the Status block in `CLAUDE.md`.

- **[VERIFIED-LIVE]** — exercised in the 2026-07 test run: certificate import into ACM, offline image builds on the hardened host, ALB + access logs, DB bootstrap + app-user auth, RDS TLS (verify-full via the OS trust store), and the endpoint-SG reachability model.
- **[NEEDS TEST-RUN CONFIRMATION]** — doc-verified against the scripts and templates but **not yet exercised end to end**: gateway steady state + login, secret rotation (DB app credential, Okta, Grafana, portal), Grafana Okta login, the activity archive, alarm firing, restore, teardown, and the entire **download portal** (stack `04`: live Okta round-trip with the groups claim, real-size streamed downloads, audit-log wiring). The asynchronous db-secret rotation event shape is the standing example (`CLAUDE.md`).

Treat a whole runbook's tag as the ceiling; individual steps call out anything more specific.

Accepted risks (surfaced up front)

Two deliberate, SSP-scoped decisions affect operations and are stated here so they are not discovered mid-incident:

- **C2 — plaintext OTLP hop.** The gateway→collector telemetry hop is plaintext but SG-scoped (the TLS recipe is documented on the collector task). Restarts and image rolls of the collector do not change this.
- **C9 — S3 Object Lock deferred.** The activity archive and ALB-log buckets rely on `DeletionPolicy: Retain` + bucket lifecycle, not Object Lock. A privileged operator can still delete archived objects; there is no WORM guarantee.

1. ALB TLS certificate rotation

Trigger / Frequency: The `${NAME_PREFIX}-certificate-expiry` CloudWatch alarm fires (`AWS/CertificateManager DaysToExpiry ≤ CERT_EXPIRY_ALARM_DAYS`, default 30) — or any unplanned re-issue (key compromise, CA change). **Imported ACM certificates do NOT auto-renew**, so this is a scheduled human task, roughly once per certificate lifetime (typically annually). Status: **[NEEDS TEST-RUN CONFIRMATION]** for the in-place replace + listener pickup; the initial import path is **[VERIFIED-LIVE]**.

Preconditions:

- Access to the enterprise CA to sign a new leaf (serverAuth EKU).
- `CERTIFICATE_ARN` and `GATEWAY_FQDN` set in `deploy.env` (the current cert's ARN — rotation replaces the certificate *in place* under the same ARN, so no stack update is needed and the ALB keeps its DNS name).
- Run the key/CSR steps on the PKI workstation; `import-enterprise-cert.sh` sources `common.sh` with `COMMON_SH_OPTIONAL_ENV=1`, so it works without a filled-in `deploy.env` (it only needs `set_env_var`, which no-ops with a warning if `deploy.env` is absent).

Steps (exact commands):

1. Generate a fresh key + CSR. The script does this under `umask 077` and removes any pre-existing key first, so the key is never briefly world-readable:

```
bash ./scripts/import-enterprise-cert.sh csr "$GATEWAY_FQDN" # writes ${GATEWAY_FQDN}.key.pem (0600) + ${GATEWAY_FQDN}.csr # SAN is exactly
DNS:${GATEWAY_FQDN} (the corporate CNAME, not the ALB name) # Key type defaults to EC P-256; append rsa2048 (or rsa3072) if the CA # only
issues RSA: ...csr "$GATEWAY_FQDN" rsa2048
```

1. Submit `${GATEWAY_FQDN}.csr` to the enterprise CA (serverAuth EKU). Collect the new leaf (`leaf.pem`) and the CA chain (`chain.pem`, intermediates first, root last).
2. **Publish the new SHA-256 fingerprint to developers BEFORE cutting over.** Rotation re-triggers Claude Code's first-connect trust prompt; developers who pinned the old fingerprint must be told the new one. The import step prints it, but you can print it ahead of the cutover from the leaf:

```
bash openssl x509 -in leaf.pem -noout -fingerprint -sha256
```

1. Replace the certificate **in place** under the existing ARN — the ALB listener picks up the new material with **no stack update**:

```
bash ./scripts/import-enterprise-cert.sh import "$GATEWAY_FQDN" \ leaf.pem "${GATEWAY_FQDN}.key.pem" chain.pem \ --certificate-arn
"$CERTIFICATE_ARN"
```

The script re-validates SAN, serverAuth EKU, and key↔cert match before importing; on success it re-prints the fingerprint and the new expiry date, and persists `CERTIFICATE_ARN` back into `deploy.env` (unchanged on an in-place replace).

Verification:

- `scripts/verify-gateway.sh` — step 2/3 fetches the served cert, cross-checks its SHA-256 against the ACM-imported cert (when AWS creds are available), and prints the fingerprint to publish. A mismatch flags TLS interception (ZIA inspection), not a rotation success.
- Confirm the alarm clears: after ACM re-computes `DaysToExpiry` (daily, Period: 86400), `${NAME_PREFIX}-certificate-expiry` returns to `OK`.

Rollback / recovery: Re-import the previous leaf/key/chain under the same `--certificate-arn` (keep the outgoing material until the new cert is confirmed live). Because the ARN is stable, rollback is another in-place `import`; the ALB never changes DNS name or listener config.

Notes & pitfalls:

- **Do not** delete-and-recreate the ACM certificate or change the listener's `CertificateArn` in the template to rotate — that risks an ALB/listener update. In-place replace under the same ARN is the only sanctioned path.
 - The SAN must be `DNS:${GATEWAY_FQDN}` (the corporate CNAME), never the `*.elb.amazonaws.com` name.
 - **Test-account variant:** the self-signed shortcut (a leaf that is its own trust anchor, imported with no chain) lives in [test-run-runbook.md §1 "Test-account shortcut — self-signed ALB cert"](#). Production rotation always uses an enterprise-CA leaf + chain as above; the self-signed path is a cross-reference only.
-

2. Okta client-secret rotation (gateway, Grafana & portal)

Trigger / Frequency: Okta client secret expiry/rotation policy, suspected exposure, or an Okta app re-key. The gateway OIDC secret, the Grafana SSO secret, and (when stack `04` is deployed) the download portal's OIDC secret all ride the same `put_secret_and_roll` helper (hidden prompt → mode-600 `file://` write → forced ECS new-deployment). Status: **[NEEDS TEST-RUN CONFIRMATION]** (steady-state login not yet exercised).

Preconditions:

- **Coordinate with the Okta admin first.** Generate the new client secret in Okta and have its value in hand *before* running the script. Okta apps can hold two secrets during an overlap window — ask the admin to add the new secret without removing the old one, so there is no outage between "secret written to Secrets Manager" and "old secret retired in Okta".
- `GATEWAY_STACK_NAME` deployed (gateway secret); `OBS_STACK_NAME` deployed (Grafana secret); `PORTAL_STACK_NAME` deployed (portal secret).

Steps (exact commands):

- **Gateway OIDC secret** — reads `OktaClientSecretArn`, `ClusterName`, `ServiceName` from the gateway stack and rolls the gateway service:

```
bash ./scripts/set-okta-secret.sh # prompts "Okta client secret (input hidden):" — paste the NEW value
```

- **Grafana SSO secret** — reads `GrafanaOidcSecretArn` + `GrafanaServiceName` from the observability stack and `ClusterName` from the gateway stack (shared cluster), and rolls the Grafana service:

```
bash ./scripts/set-grafana-oidc-secret.sh # prompts "Okta client secret for Grafana (input hidden):" — paste the NEW value
```

If Grafana reuses the gateway's Okta app, this is the *same* secret value as `set-okta-secret.sh`.

- **Portal OIDC secret** (when stack 04 is deployed) — reads `PortalOidcSecretArn` + `PortalServiceName` from the portal stack and `ClusterName` from the gateway stack (shared cluster), and rolls the portal service:

```
bash ./scripts/set-portal-oidc-secret.sh # prompts hidden — paste the NEW value
```

If the portal reuses the gateway's Okta app (a documented option — `scripts/deploy.env.example`), this is the *same* secret value as `set-okta-secret.sh`; rotate **both** consumers in the same window.

Verification:

- Watch the roll to stable. The `put_secret_and_roll` helper echoes the exact `aws ecs wait services-stable ...` command (with the resolved cluster/service) at the end of its run — copy and run that line.
- Gateway: `scripts/verify-gateway.sh` step 3/3 — the OAuth endpoints respond and issue a device `user_code`.
- Grafana: sign in at `https://${GATEWAY_FQDN}/grafana` via Okta.
- Portal: sign in at `https://${GATEWAY_FQDN}/portal` via Okta and reach the Team/Cost-Center page.

Rollback / recovery: If login breaks after the roll, re-run the same script and paste the **previous** secret value (Okta still honours it during the overlap window). Then investigate before retrying.

Failure mode — rolling before Okta has the new value: If you write a new secret to Secrets Manager and roll the service while Okta still expects the old one, the OAuth code exchange fails and logins break for everyone (Grafana with the login form disabled = total lockout). Always confirm the new secret is active in Okta first; keep the old secret valid until the new tasks are stable.

Notes & pitfalls:

- Never pass the secret on a command line — the scripts prompt for it (hidden) and write via a mode-600 `file://` temp file (`.claude/rules/security.md`).
- The template resources are placeholders (`REPLACE-ME-...`). Do **not** "rotate" by editing the template `SecretString` — that clobbers the live value on the next deploy (`.claude/rules/cloudformation.md`).

3. Database app-credential rotation

Trigger / Frequency: Normally **automatic** — the db-admin rotation Lambda (`${NAME_PREFIX}-db-rotation`) runs on the `APP_SECRET_ROTATION_DAYS` cadence (`deploy.env`, default 90; passed through as the `AppSecretRotationDays` template parameter). Manual triggers: suspected credential exposure, or recovery after

a half-completed rotation. Status: **[NEEDS TEST-RUN CONFIRMATION]** — rotation is not yet exercised in steady state, and the asynchronous rotation event shape is doc-verified only (`CLAUDE.md`).

Design (how it works — describe, don't guess): The secret `${NAME_PREFIX}/db-app-user` alternates between two Postgres LOGIN users, `gateway_app` and `gateway_app_clone`, both of which assume the NOLOGIN owner role `gateway_owner` at login. Each rotation flips `AWSCURRENT` to the *other* user with a fresh password, so the **previous credential stays valid until the next rotation** — there is no window where a running task holds a dead credential. The four standard Secrets Manager steps (`docker/db-admin/app.py:rotate_handler`, semantics fixed by `tests/lambda/test_rotation.py`):

1. `createSecret` — put the pending value (other user + random password) at `AWSPENDING` (idempotent on retry).
2. `setSecret` — `ALTER ROLE <pending user> WITH PASSWORD ...` as the RDS master.
3. `testSecret` — connect as the pending user and `SELECT 1`.
4. `finishSecret` — move `AWSCURRENT` to the new version (idempotent), then `forceNewDeployment` on the gateway service so new tasks fetch the new credential.

Preconditions: `DB_STACK_NAME` + `GATEWAY_STACK_NAME` deployed; the db-admin Lambda image is current.

Steps (exact commands):

- **Confirm the last rotation succeeded:**

```
```bash # Rotation metadata: last-rotated date, schedule, and which versions hold # which stages (expect one AWSCURRENT, and AWSPREVIOUS = the still-
valid # prior credential). aws secretsmanager describe-secret --region "$AWS_REGION" --secret-id "${NAME_PREFIX}/db-app-user" \ --query
'{LastRotated:LastRotatedDate, RotationEnabled:RotationEnabled, Stages:VersionIdsToStages}'
```

```
Rotation Lambda log group — look for "rotation finished; service roll requested" aws logs tail "/aws/lambda/${NAME_PREFIX}-db-rotation" --region
"$AWS_REGION" --since 100d ```
```

- **Trigger a manual rotation:**

```
bash aws secretsmanager rotate-secret --region "$AWS_REGION" \ --secret-id "${NAME_PREFIX}/db-app-user"
```

(Rotation is asynchronous; the CLI returns immediately. Watch the log group above for the four steps.)

*Verification:*

- `describe-secret` shows a newer `LastRotatedDate` and `AWSCURRENT` on a new version id; `AWSPREVIOUS` points at the prior version.
- The `${NAME_PREFIX}-db-rotation-errors` alarm stays `OK` (threshold:  $\geq 3$  Lambda errors/hour — it tolerates the expected single inactive-image error per scheduled rotation).
- Gateway tasks are stable after the `finishSecret` roll ( `aws ecs wait services-stable ... --services $NAME_PREFIX-gateway` ).

*Rollback / recovery — half-completed rotation:* The design is retry-first: Secrets Manager retries a failed step, `finishSecret`'s label move is idempotent, and the prior credential remains valid, so a stuck rotation does **not** break running tasks — they keep working on the current credential.

- If a rotation is wedged, inspect `/aws/lambda/$NAME_PREFIX-db-rotation` for the failing step, fix the cause (common: the image Lambda is `Inactive` and the first invoke fails while Lambda re-optimizes — Secrets Manager's retries then complete it), and re-run `rotate-secret` to re-drive it.
- To abandon an in-flight `AWSPENDING` version without applying it, remove the `AWSPENDING` stage from that version (`update-secret-version-stage --remove-from-version-id ...`); the live `AWSCURRENT` credential is untouched.
- **Do not** hand-edit `${NAME_PREFIX}/db-app-user` — it is Lambda-managed (least-privilege AC-6 design). Hand-editing desyncs the secret from the Postgres role passwords.

*Notes & pitfalls:*

- The gateway **never** uses the RDS master credential; the master secret is break-glass only (see runbook 7). Rotation `ALTER ROLE`s run as master inside the Lambda, not from any task.
- Flag per `.claude/rules/process.md`: the async rotation/EventBridge event shape is **doc-verified only until the test run confirms it**.

---

## 4. RDS CA bundle refresh

*Trigger / Frequency:* AWS rotates the RDS server CA (e.g. the `rds-ca-rsa2048-g1` family — the instance's `CACertificateIdentifier`), or you must move to a newer CA before an AWS-announced expiry. Rare (multi-year). Status: **[NEEDS TEST-RUN CONFIRMATION]** for the CA-change cutover; the image build + baked-bundle mechanism is **[VERIFIED-LIVE]**.

*Why this is an image rebuild, not a config flip:* both the gateway and the db-admin Lambda connect with `sslmode=verify-full`, and the driver trusts the **OS/container trust store**, not `sslrootcert=` (proven in the test run — the driver ignores `sslrootcert=`). The RDS CA bundle is fetched at build time (`RDS_CA_BUNDLE_URL`, default the GovCloud truststore) and **baked into both images** (`docker/rds-ca-bundle.pem`, `docker/db-admin/rds-ca-bundle.pem`). A CA change therefore means: restage the bundle → rebuild **both** images with a **bumped immutable tag** → stack update that rolls the services and re-points the Lambda images.

*Preconditions:* Build host with Docker + egress to the RDS truststore (or a mirrored bundle), `KMS_KEY_ARN` set (CMK-encrypted ECR), and — if the CA identifier itself changes on the instance — a maintenance window (modifying `CACertificateIdentifier` on the DB may require a reboot).

*Steps (exact commands):*

1. **Rebuild the gateway image with a bumped tag** (tags are IMMUTABLE — a same-tag rebuild cannot be pushed). The build script re-fetches the RDS CA bundle every run:

```
bash # bump the tag so the new bundle ships under a new immutable URI IMAGE_TAG="${CLAUDE_VERSION}-ca$(date +%Y%m%d)" ./scripts/build-and-push-image.sh # persists IMAGE_URI back into deploy.env
```

### 1. Rebuild the db-admin Lambda image with a bumped tag:

```
bash DBADMIN_VERSION="1.0.1" ./scripts/build-and-push-dbadmin.sh # persists DBADMIN_IMAGE back into deploy.env
```

(Optionally override `RDS_CA_BUNDLE_URL` on both builds to pin a specific bundle for a controlled network.)

1. **Deploy the gateway stack** so the task definition and both Lambdas pick up the new image URIs and the service rolls (images **before** the stack update that expects them — `.claude/rules/scripts.md`):

```
bash ./scripts/deploy-gateway.sh
```

1. **Only if the instance CA identifier changes:** update `CACertificateIdentifier` on the RDS instance in `01-database.yaml` to the new CA and `./scripts/deploy-database.sh`. This is a property modification, not a replacement — verify it is not flagged as `Update:Replace` before applying (the stack policy denies replacement of `Database`).

*Verification:*

- Gateway + db-admin connect with `verify-full` against the new CA: `scripts/verify-gateway.sh` passes end to end, and a manual `aws secretsmanager rotate-secret --secret-id "$NAME_PREFIX/db-app-user"` completes its `testSecret` step (proves the rebuilt db-admin image trusts the new CA).
- `aws ecs wait services-stable ... --services $NAME_PREFIX-gateway`.

*Rollback / recovery:* Redeploy with the previous `IMAGE_URI` / `DBADMIN_IMAGE` (both still exist under their old immutable tags in ECR) via `deploy.env` + `deploy-gateway.sh`. Because AWS RDS CA changes are additive (old + new CA trusted during the transition window), the old images keep validating until the old CA is retired.

*Notes & pitfalls:*

- Rebuild **both** images — a refreshed gateway with a stale db-admin image (or vice versa) leaves one side unable to validate `verify-full` after the CA fully cuts over.
  - Do not attempt to fix a CA change by setting `sslrootcert=` — the driver ignores it. The only lever is the baked-in bundle.
-

## 5. Claude Code release update

*Trigger / Frequency:* A new pinned Claude Code release you want to distribute (security fix, feature, or a gateway-required minimum bump). Status: **[NEEDS TEST-RUN CONFIRMATION]** for the full mirror→build→distribute chain; offline image builds are **[VERIFIED-LIVE]**.

*Preconditions:*

- An egress host that can reach `downloads.claude.ai` (the mirror step). The laptops and the container build need **no** egress.
- `ANTHROPIC_GPG_KEY` set to Anthropic's release-signing public key so the manifest signature is verified. Verification **fails closed**: `ALLOW_UNVERIFIED_MANIFEST=1` is the only (deliberate, named) escape hatch and must not be the default (`.claude/rules/security.md`).

*Steps (exact commands):*

1. **Mirror the release** (linux-x64 for the image + win32-x64 for laptops):

```
bash ANTHROPIC_GPG_KEY=/path/to/anthropic-release-key.asc \ ./client/mirror-claude-release.sh 2.1.208 # verifies the GPG-signed manifest + per-binary SHA-256, writes # mirror/2.1.208/{claude,claude.exe,CHECKSUMS.txt}
```

1. **Rebuild the gateway image** (it embeds the linux binary). Set `CLAUDE_VERSION` to the new release in `deploy.env` first, then stage the verified binary and build:

```
bash cp mirror/2.1.208/claude docker/claude # deploy.env: export CLAUDE_VERSION="2.1.208" ./scripts/build-and-push-image.sh # tags the image 2.1.208, persists IMAGE_URI ./scripts/deploy-gateway.sh # rolls the gateway service onto it
```

1. **Publish to the download portal** (when stack 04 is deployed) — this is how developers self-serve the new version. Reuses the verified mirror output; uploads `claude.exe`, `manifest.json`, `CHECKSUMS.txt`, and the installer to the portal's CMK-encrypted artifacts bucket, then pins the portal to the new version:

```
bash ./scripts/publish-portal-release.sh 2.1.208 # deploy.env: export PORTAL_RELEASE_VERSION="2.1.208" (empty = CLAUDE_VERSION) ./scripts/deploy-download-portal.sh # only needed when the pinned version changes
```

1. **Distribute the Windows client (share/MDM route)**. Stage `mirror/2.1.208/claude.exe` + `CHECKSUMS.txt` on the file share and install non-elevated per developer:

```
powershell .\client\Install-ClaudeCode.ps1 -BinaryPath \\fileserver\software\claude\2.1.208\claude.exe -Sha256 <win32-x64 checksum from CHECKSUMS.txt> -GatewayUrl https://<GATEWAY_FQDN> -DisableUpdates
```

1. **Forcing the upgrade (optional)**. To make the CLI refuse to start below the new version, bump `-RequiredMinimumVersion` in the managed-settings push (writes `requiredMinimumVersion` into `managed-settings.json`; default floor is `2.1.195`, the gateway's minimum):

```
powershell .\client\Install-ClaudeCode.ps1 -SettingsOnly ` -GatewayUrl https://<GATEWAY_FQDN> -DisableUpdates ` -RequiredMinimumVersion 2.1.208
```

#### Verification:

- Mirror step: the script prints `checksum OK` per platform and `manifest signature OK`. A SHA-256 or signature mismatch aborts (non-zero exit) and removes the bad file.
- Gateway: `aws ecs wait services-stable ... --services $NAME_PREFIX-gateway`, then `scripts/verify-gateway.sh`.
- Client: `claude --version` reports the new version; below-floor binaries refuse to start when `requiredMinimumVersion` is raised.
- Portal (if published): download a ZIP from `https://${GATEWAY_FQDN}/portal` and confirm it contains the new `claude.exe` and that the generated `install.cmd` carries the new version's SHA-256; the download appears in the portal audit log group (`/claude/${NAME_PREFIX}/portal-audit`).

*Rollback / recovery:* Redeploy the previous `IMAGE_URI` (old immutable tag still in ECR) via `deploy.env` + `deploy-gateway.sh`; re-push managed settings with the prior `-RequiredMinimumVersion` if you raised the floor. Portal: set `PORTAL_RELEASE_VERSION` back to the prior version and re-run `deploy-download-portal.sh` (earlier `releases/<version>/` prefixes stay in the artifacts bucket). Keep the prior `mirror/<version>/` directory until the new release is confirmed across the fleet.

#### Notes & pitfalls:

- Update lockdown (`-DisableUpdates` → `DISABLE_UPDATES=1` + `DISABLE_AUTOUPDATER=1`) is what keeps users on the distributed version — do not drop it, or clients will self-update off the pinned build.
- SYSTEM-context (Intune/SCCM device) pushes must use `-SettingsOnly`; the binary install must run in **user** context (the installer throws otherwise).
- Never bypass GPG verification as a matter of routine; `ALLOW_UNVERIFIED_MANIFEST=1` is for a deliberately air-gapped one-off only.

---

## 6. Gateway / Grafana / collector image & stack updates

*Trigger / Frequency:* Any container change (Dockerfile fix, Grafana provisioning, a new ADOT collector release, task CPU/memory tuning, or a parameter change). Status: **[NEEDS TEST-RUN CONFIRMATION]** for steady-state rolls of Grafana/collector; gateway image builds + deploys are **[VERIFIED-LIVE]**.

*Preconditions:* Build host with Docker; `KMS_KEY_ARN` set; the relevant stack already deployed. **Rebuild and push the image BEFORE the stack update that expects it** (`.claude/rules/scripts.md`).

*Steps (exact commands):* bump the immutable tag → build/push → deploy.

- **Gateway:** `IMAGE_TAG=<new> ./scripts/build-and-push-image.sh → ./scripts/deploy-gateway.sh`.
- **Grafana:** `GRAFANA_IMAGE_TAG=<new> ./scripts/build-and-push-grafana.sh → ./scripts/deploy-observability.sh`.

- **ADOT collector:** `ADOT_VERSION=<vX.Y.Z> ./scripts/mirror-collector.sh` (mirrors + pins `COLLECTOR_IMAGE` by digest) → `./scripts/deploy-observability.sh`.
- **db-admin Lambda:** `DBADMIN_VERSION=<new> ./scripts/build-and-push-dbadmin.sh` → `./scripts/deploy-gateway.sh`.
- **Download portal:** `PORTAL_VERSION=<new> ./scripts/build-and-push-portal.sh` → `./scripts/deploy-download-portal.sh`.

Each build script persists its new URI/tag into `deploy.env`, so the matching `deploy-*.sh` picks it up with no copy-paste.

*Update-safety invariants (must hold on every stack update — `.claude/rules/cloudformation.md`):*

- **The ALB and RDS instance must never be replaced by a routine update.** Both are protected three ways — deletion protection, fixed physical names, and a **stack policy** (set by `deploy-gateway.sh` / `deploy-database.sh`) denying `Update:Replace` / `Update>Delete` on `LoadBalancer` / `Database`. Do not remove any layer. `deploy-gateway.sh` deploys with `--disable-rollback` by default (`CFN_DISABLE_ROLLBACK=true`) precisely so a failed create keeps the protected ALB rather than attempting an impossible rollback delete.
- **Cross-stack exports are locked while imported.** 01 exports the CMK, DB endpoint, master-secret ARN, and client SG to 02; 02 exports SGs, the listener, and the cluster to 03. You cannot change an exported value in place while a downstream stack imports it — encryption-at-rest and resource names are day-one decisions.
- **Placeholder `SecretString` resources must not be touched.** `OktaClientSecret`, `Grafana0idcClientSecret`, `Portal0idcClientSecret`, and `DbAppUserSecret` hold placeholder/managed values; editing the `SecretString` literal (or the resource Name/Description) re-applies the placeholder and clobbers the live secret. Rotate via the scripts (runbooks 2–3), never the template.
- `TaskCpu` / `TaskMemory` **must stay a valid Fargate pairing** — the template's `Rules` section asserts this; an invalid combo fails deploy with an opaque error. Change them via `TASK_CPU` / `TASK_MEMORY` in `deploy.env`.

*Verification:* `aws ecs wait services-stable` for the affected service (`$NAME_PREFIX-gateway`, `$NAME_PREFIX-grafana`, `$NAME_PREFIX-otel`, or `$NAME_PREFIX-portal`); `scripts/verify-gateway.sh` for the gateway; Grafana login + dashboards for Grafana. Confirm the CloudFormation events show `UPDATE_COMPLETE` with **no** replacement of `LoadBalancer` or `Database`.

*Rollback / recovery:* Redeploy the prior image URI/tag from `deploy.env` (old immutable tags persist in ECR). For a parameter regression, re-run the deploy script with the previous `deploy.env` values. If a gateway create/update lands in `*_FAILED` with `--disable-rollback`, fix the cause and re-run `deploy-gateway.sh` — the deploy **continues** from where it failed rather than tearing down the protected ALB.

*Notes & pitfalls:* A same-tag rebuild cannot be pushed (immutable repos) and an unchanged image URI leaves the service/Lambda on old code — always bump the tag.

---

## 7. Secrets inventory & break-glass

*Trigger / Frequency:* Reference during audits, incident response, or before any secret change. Status: **[NEEDS TEST-RUN CONFIRMATION]** for the break-glass master path.

Secrets inventory (all CMK-encrypted with the CMK from 01):

Secret (Name)	Stack	Rotation	How to rotate
RDS master \${NAME_PREFIX} DB (Database.MasterUserSecret)	01	<b>Automatic</b> , RDS-managed, every 7 days	RDS-managed; break-glass use → force-rotate (below)
\${NAME_PREFIX}/db-app-user	02	<b>Automatic</b> , alternating-user Lambda, AppSecretRotationDays (default 90)	Runbook 3
\${NAME_PREFIX}/oidc-client-secret (Okta)	02	<b>Manual</b>	scripts/set-okta-secret.sh (runbook 2)
\${NAME_PREFIX}/jwt-secret (session signing)	02	<b>Not rotated automatically</b> (GenerateSecretString at create)	Manual — see below
\${NAME_PREFIX}/grafana-oidc-client- secret	03	<b>Manual</b>	scripts/set-grafana-oidc-secret.sh (runbook 2)
\${NAME_PREFIX}/grafana-admin-password	03	<b>Not rotated</b> (break-glass; login form disabled)	Regenerate manually (below)
\${NAME_PREFIX}/portal-oidc-client- secret	04	<b>Manual</b>	scripts/set-portal-oidc-secret.sh (runbook 2)
\${NAME_PREFIX}/portal-session-secret (cookie signing)	04	<b>Not rotated automatically</b> (GenerateSecretString at create)	Same file-based pattern as the JWT secret (below), then roll \$NAME_PREFIX-portal; rotation invalidates portal sessions (users just re-login)

*Gateway JWT secret (manual rotation).* The template describes rotation as "prepend new value, roll, remove old." Whether the gateway honours two overlapping signing keys is **doc-verified only — needs test-run confirmation**; until confirmed, treat a JWT rotation as session-invalidating (active sessions below `SessionTtlHours`, default 1h, are dropped and users re-login). Rotate by writing a new value via the same safe pattern the helper uses, then forcing a roll:

```
JWT_ARN=$(aws cloudformation describe-stack-resources --region "$AWS_REGION" \
 --stack-name "$GATEWAY_STACK_NAME" --logical-resource-id JwtSecret \
 --query 'StackResources[0].PhysicalResourceId' --output text)
generate + write without ever putting the value on argv:
f=$(mktemp); chmod 600 "$f"; aws secretsmanager get-random-password \
 --region "$AWS_REGION" --password-length 48 --exclude-punctuation \
 --query RandomPassword --output text > "$f"
aws secretsmanager put-secret-value --region "$AWS_REGION" \
 --secret-id "$JWT_ARN" --secret-string "file://$f"; rm -f "$f"
aws ecs update-service --region "$AWS_REGION" \
 --cluster "$(aws cloudformation describe-stacks --region "$AWS_REGION" \
 --stack-name "$GATEWAY_STACK_NAME" \
```

```
--query "Stacks[0].Outputs[?OutputKey=='ClusterName'].OutputValue" --output text)" \
--service "$NAME_PREFIX-gateway" --force-new-deployment
```

*Grafana admin password (break-glass regenerate).* Same file-based pattern against `${NAME_PREFIX}/grafana-admin-password`, then roll `$NAME_PREFIX-grafana`. The login form stays disabled unless `GRAFANA_DISABLE_LOGIN_FORM=false`; day-to-day access is Okta SSO.

*Break-glass — RDS master secret.* The master credential is **break-glass ONLY**; no task ever injects it (the gateway uses `${NAME_PREFIX}/db-app-user`). Use it only for direct DBA access during an incident.

*Steps:*

1. Read it without echoing to the terminal history:

```
bash MASTER_ARN=$(aws cloudformation describe-stacks --region "$AWS_REGION" \ --stack-name "$DB_STACK_NAME" \ --query "Stacks[0].Outputs[?OutputKey=='DBMasterSecretArn'].OutputValue" --output text) # inspect keys interactively; avoid persisting the value anywhere
aws secretsmanager get-secret-value --region "$AWS_REGION" \ --secret-id "$MASTER_ARN" --query SecretString --output text
```

1. Perform the minimum necessary DBA action, from a host that can reach the DB (in-VPC operator host; the DB is `PubliclyAccessible: false`).

2. **Immediately afterwards — rotate the master back** so the just-exposed value is retired, and **record the use** (who/when/why) for the audit trail:

```
bash aws secretsmanager rotate-secret --region "$AWS_REGION" --secret-id "$MASTER_ARN"
```

*Verification:* `describe-secret` on the master shows a fresh `LastRotatedDate` after the forced rotation; the gateway is unaffected (it never used the master).

*Notes & pitfalls:* Treat the activity-log stream as highly sensitive (bash commands, tool inputs, file paths per user) — opt-in, IAM-only, CMK-encrypted, SIEM-flagged; never widen its access surface. Never put any secret value on a command line (`--secret-string <value>` leaks via `ps /proc`); always the `mode-600 file://` pattern above.

---

## 8. Backup & restore

*Trigger / Frequency:* Reference for DR planning; act on data-loss/corruption or before a risky change. Status: **[NEEDS TEST-RUN CONFIRMATION]** for the restore path.

*Posture:*

- **RDS automated backups** — `BackupRetentionPeriod = BackupRetentionDays` (01-database.yaml, default 14, max 35). Daily automated snapshots + PITR within the window. `DeletionPolicy: Snapshot / UpdateReplacePolicy: Snapshot` → a stack delete/replace takes a **final snapshot** rather than destroying data. `DeletionProtection: true` blocks accidental instance deletion.
- **ALB access logs** — `AlbLogsBucket` (SSE-S3; ELB delivery does not support KMS — this is the one documented CMK exception), `DeletionPolicy: Retain`, lifecycle expiry at `AlbLogRetentionDays` (default 90).
- **Activity archive** — `ActivityArchiveBucket` (CMK-encrypted, `DeletionPolicy: Retain`), lifecycle expiry at `ActivityArchiveRetentionDays` (default 731 ≈ 2y). The CloudWatch window group `/claude/${NAME_PREFIX}/activity` retains `ActivityLogWindowDays` (default 14) before the Firehose→S3 chain is the durable copy. Idle (no cost) until the gateway sets `ForwardActivityLogs=true`.
- **AMP** — the workspace is `DeletionPolicy: Retain` so a routine 03 recreate does not destroy metrics history.

*Take an on-demand snapshot (before risky changes):*

```
aws rds create-db-snapshot --region "$AWS_REGION" \
--db-instance-identifier "$NAME_PREFIX-store" \
--db-snapshot-identifier "$NAME_PREFIX-store-preop-$(date +%Y%m%d%H%M)"
```

*Restore* — *understand the blast radius first*. **A replaced RDS instance is an EMPTY database, not a restore** (`.claude/rules/cloudformation.md`), and the DB endpoint is a cross-stack export imported by 02, which is **locked while imported**. You cannot restore in place by pointing the stack at a snapshot without disturbing that export. Restore is therefore effectively a **teardown + restore**, not an update:

1. Restore the snapshot to a **new** instance out-of-band to validate the data (`aws rds restore-db-instance-from-db-snapshot ...`), or use PITR.
2. To make the restored data the live store, the sanctioned path is to bring the database stack back from the snapshot (RDS snapshot-based restore) with the same `DBInstanceIdentifier`/exports so 02 re-imports the endpoint — which, given the export lock and deletion protection, means an orchestrated teardown of 02 first, restore of 01 from the snapshot, then redeploy of 02/03. Plan this as a maintenance-window operation, not a routine update.

*Verification:* `aws rds describe-db-snapshots` shows the expected automated + manual snapshots; a test restore to a scratch instance connects and shows the expected schema/rows. After a real restore, `scripts/verify-gateway.sh` passes and the gateway serves logins.

*Rollback / recovery:* Snapshots are immutable point-in-time copies — a failed restore attempt is retried against another snapshot; the source snapshots are unaffected. Keep the pre-op on-demand snapshot until the operation is confirmed.

*Notes & pitfalls:* Per C9 (accepted risk), neither S3 bucket uses Object Lock — archived logs are deletable by a privileged operator. If tamper-evidence becomes a requirement, revisit C9 in the security review.

---

## 9. Alarm response

*Trigger / Frequency:* On alarm (routed to `ALARM_SNS_TOPIC_ARN` when set — otherwise the alarms exist but have no action). Status: **[NEEDS TEST-RUN CONFIRMATION]** — alarms are defined but have not fired in the test run.

*Alarms defined in the templates:*

1. `${NAME_PREFIX}-certificate-expiry` (`02-gateway.yaml`) — `AWS/CertificateManager` `DaysToExpiry ≤ CERT_EXPIRY_ALARM_DAYS` (default 30).  
*Response:* the imported cert is approaching expiry and will **not** auto-renew → execute **runbook 1** (re-issue from the enterprise CA, publish the new fingerprint, in-place `import --certificate-arn`). Alarm clears once ACM recomputes `DaysToExpiry` (daily).
2. `${NAME_PREFIX}-db-rotation-errors` (`02-gateway.yaml`) — `AWS/Lambda` `Errors ≥ 3` in one hour on `${NAME_PREFIX}-db-rotation`. *Response:* rotation is erroring repeatedly (running tasks are still fine on the current credential, but the rotation SLA is at risk). Inspect `/aws/lambda/${NAME_PREFIX}-db-rotation`, identify the failing step, and follow **runbook 3** recovery. The threshold intentionally tolerates the single expected Inactive-image error per scheduled rotation.

*No other CloudWatch alarms are defined in the templates* (`03-observability` has Cloud Map health checks and target-group health thresholds, but no `AWS::CloudWatch::Alarm` resources). Operational surfaces to watch manually: ECS service events / `services-stable`, and the gateway/Grafana/collector log groups.

*General verification after responding:* confirm the alarm returns to `OK` (`aws cloudwatch describe-alarms --alarm-names <name> --query 'MetricAlarms[0].StateValue'`).

*Known landing-zone gotcha — ALB access-log AccessDenied.* If ALB access-log enablement fails `AccessDenied` on a bucket policy that is correct, **suspect a landing-zone auto-remediation rewriting the ALB's log config before suspecting the bucket policy** (test-run lesson). The bucket policy already grants both ELB delivery principals; the transient post-deploy log-enable variant was removed once the environment auto-remediation was exempted. Get the auto-remediation exempted for this ALB rather than re-editing the policy.

*Rollback / recovery:* Alarm response is corrective, not stateful — there is nothing to roll back beyond the underlying runbook's own recovery.

---

## 10. Teardown

*Trigger / Frequency:* Decommissioning the deployment (test account cleanup or end of life). Rare and deliberate. Status: **[NEEDS TEST-RUN CONFIRMATION]**.

*Order is the reverse of deploy:* `04` and `03` → `02` → `01`. The portal (`04`) and observability (`03`) stacks both import from `02` and are independent of each other — delete them (in either order, or in parallel) before the gateway. There is intentionally **no teardown script**; delete stacks explicitly so each destructive step is a conscious act. Downstream stacks import upstream exports, so an out-of-order delete fails on the export lock.

*Preconditions & the protection layers you must clear first:*

- **RDS** `DeletionProtection: true` blocks deleting the DB stack — disable it first (a stack update setting it false, or the console), and expect a **final snapshot** (`DeletionPolicy: Snapshot`).
- **ALB** `deletion_protection.enabled: true` blocks deleting the gateway stack — disable it first.
- The **stack policies** set by the deploy scripts deny `Update:Replace` / `Update:Delete` on `LoadBalancer` / `Database` during *updates*; they do not block `delete-stack`, but the deletion-protection flags above do. Clear those flags before deleting.

*Steps (exact commands):*

```
4) Download portal (if deployed)
aws cloudformation delete-stack --region "$AWS_REGION" --stack-name "$PORTAL_STACK_NAME"
aws cloudformation wait stack-delete-complete --region "$AWS_REGION" --stack-name "$PORTAL_STACK_NAME"

3) Observability
aws cloudformation delete-stack --region "$AWS_REGION" --stack-name "$OBS_STACK_NAME"
aws cloudformation wait stack-delete-complete --region "$AWS_REGION" --stack-name "$OBS_STACK_NAME"

2) Gateway (disable ALB deletion protection first, then delete)
aws cloudformation delete-stack --region "$AWS_REGION" --stack-name "$GATEWAY_STACK_NAME"
aws cloudformation wait stack-delete-complete --region "$AWS_REGION" --stack-name "$GATEWAY_STACK_NAME"

1) Database (disable RDS deletion protection first; a final snapshot is taken)
aws cloudformation delete-stack --region "$AWS_REGION" --stack-name "$DB_STACK_NAME"
aws cloudformation wait stack-delete-complete --region "$AWS_REGION" --stack-name "$DB_STACK_NAME"
```

*What survives deletion (verified against the templates' `DeletionPolicy`):*

- **KMS CMK** (`KmsKey`, 01) — `Retain`. Everything at rest was encrypted with it; retained so retained data stays readable. Schedule key deletion manually only after all encrypted artifacts are gone.
- **ALB access-logs bucket** (`AlbLogsBucket`, 02) — `Retain`.
- **Activity-archive bucket** (`ActivityArchiveBucket`, 03) — `Retain` (CMK-encrypted).
- **AMP workspace** (`Workspace`, 03) — `Retain`.
- **Portal artifacts bucket** (`ArtifactsBucket`, 04) — `Retain` (CMK-encrypted; holds the published release binaries).
- **Every CloudWatch log group, in every stack** — `Retain` (operator decision 2026-07-18: logs outlive stacks; a `cfn-guard` gate `log_groups_survive_tearardown` enforces it). This covers the ECS task groups (`gateway/collector/grafana/portal`), the activity window, the portal download-audit group, the RDS `postgresql/pgaudit` export group (`/aws/rds/instance/${NAME_PREFIX}-store/postgresql`, pre-created by 01 with the CMK), and the db-admin Lambda groups (pre-created by 02 with the CMK). Deleting a retained group afterwards is a deliberate manual act. Note the redeploy consequence: the groups carry fixed names, so a later **re-create collides** with the retained groups and the new stack fails — export what you need, then

delete them first (the test-run runbook §0 has the command list). The adopted groups (RDS postgresql, Lambda) additionally collide on the *first* deploy of this change into an account where the services already auto-created them.

- **RDS instance** (Database, 01) — `DeletionPolicy: Snapshot` → a **final snapshot** is taken; the running instance is removed. The snapshot persists.
- Everything else (ALB, ECS services/cluster, secrets, VPC endpoints, Lambdas, Firehose) is **deleted** with its stack.

*Verification:* all three `stack-delete-complete` waits return; `aws cloudformation describe-stacks` reports the stacks gone; confirm the retained buckets, CMK, AMP workspace, and final DB snapshot still exist if you intend to keep them, or clean them up explicitly.

*Rollback / recovery:* Redeploy from scratch per the [test-run-runbook](#). Data recovery relies on the retained final RDS snapshot (restore per runbook 8) and the retained buckets.

*Notes & pitfalls:* Deleting a stack that another stack still imports from fails — always 04 and 03 → 02 → 01, waiting for each delete to complete before the next tier. Retained resources are **not** free; account for the retained buckets, CMK, AMP workspace, and snapshots after teardown.